

```
In [1]: ▶ import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from scipy.signal import find_peaks
from mpl_toolkits.mplot3d import Axes3D
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report, confusi
from sklearn.metrics import accuracy_score, precision_score, recall_score,
```

```
In [2]: ▶ #import the train and test datasets
df0 = pd.read_csv('mitbih_train.csv', header = None)
df1 = pd.read_csv('mitbih_test.csv', header = None )
```

```
In [3]: ▶ #show the number of columns in each dataframe
columns_df0 = df0.columns
print("'mitbih_train.csv: ", columns_df0)
print(' ')
print('#####')
print(' ')
columns_df1 = df1.columns
print("'mitbih_test.csv: ", columns_df1)
```

```
'mitbih_train.csv: Int64Index([ 0,  1,  2,  3,  4,  5,  6,  7,
8,  9,
```

```
...
178, 179, 180, 181, 182, 183, 184, 185, 186, 187],
dtype='int64', length=188)
```

```
#####
####
```

```
mitbih_test.csv: Int64Index([ 0,  1,  2,  3,  4,  5,  6,  7,
8,  9,
```

```
...
178, 179, 180, 181, 182, 183, 184, 185, 186, 187],
dtype='int64', length=188)
```

```
In [4]: ▶ # concatenate the dataframes to make one dataframe dfs
dfs = [df0, df1]
dfs = pd.concat(dfs, axis=0)
```

```
In [5]: dfs.head()
```

Out[5]:

	0	1	2	3	4	5	6	7	8
0	0.977941	0.926471	0.681373	0.245098	0.154412	0.191176	0.151961	0.085784	0.058824
1	0.960114	0.863248	0.461538	0.196581	0.094017	0.125356	0.099715	0.088319	0.074074
2	1.000000	0.659459	0.186486	0.070270	0.070270	0.059459	0.056757	0.043243	0.054054
3	0.925414	0.665746	0.541436	0.276243	0.196133	0.077348	0.071823	0.060773	0.066298
4	0.967136	1.000000	0.830986	0.586854	0.356808	0.248826	0.145540	0.089202	0.117371

5 rows × 188 columns



```
In [6]: dfs.info()#data shows all the information about the datafarme including nu  
#and the data types and amount of memory
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 109446 entries, 0 to 21891  
Columns: 188 entries, 0 to 187  
dtypes: float64(188)  
memory usage: 157.8 MB
```

```
In [7]: # check for missing values, we can see that there are no null values in th  
dfs.isnull().sum()
```

Out[7]:

0	0
1	0
2	0
3	0
4	0
..	
183	0
184	0
185	0
186	0
187	0

Length: 188, dtype: int64

```
In [8]: ▶ # finding out the number of unique values in each column
# we can see that in column 187 there are five unique values which stand f
#of the different ecg classes
dfs.nunique()
```

```
Out[8]: 0      26912
1      48772
2      57132
3      46919
4      44736
...
183     1128
184     1033
185       961
186       912
187         5
Length: 188, dtype: int64
```

```
In [9]: ▶ dfs.shape #shape of the merged dataframe showing 109446, rows and 188 colu
```

```
Out[9]: (109446, 188)
```

```
In [10]: ▶ #randomise the dataframe and reset index
dfs = dfs.sample(frac=1).reset_index(drop = True)
```

```
In [ ]: ▶
```

In [11]: `dfs.head(20)`

Out[11]:

	0	1	2	3	4	5	6	7	
0	1.000000	0.691667	0.166667	0.013889	0.038889	0.036111	0.005556	0.002778	0.000000
1	1.000000	0.774487	0.232346	0.143508	0.129841	0.104784	0.077449	0.091116	0.059222
2	1.000000	0.892982	0.428070	0.128070	0.182456	0.198246	0.157895	0.154386	0.156140
3	1.000000	0.831099	0.310992	0.040214	0.040214	0.053619	0.029491	0.008043	0.005366
4	1.000000	0.907692	0.712821	0.533333	0.307692	0.153846	0.123077	0.174359	0.164103
5	0.975332	0.747628	0.018975	0.011385	0.000000	0.024668	0.062619	0.096774	0.092973
6	1.000000	0.820475	0.495549	0.197329	0.000000	0.060831	0.129080	0.111276	0.091987
7	0.995614	0.894737	0.320175	0.006579	0.160088	0.175439	0.094298	0.072368	0.072368
8	0.898734	0.857595	0.496835	0.000000	0.072785	0.205696	0.300633	0.306962	0.341777
9	0.000000	0.010336	0.060724	0.142119	0.241602	0.310078	0.383721	0.454780	0.498700
10	0.990148	0.889984	0.431855	0.014778	0.050903	0.210181	0.231527	0.210181	0.215100
11	1.000000	0.810127	0.396624	0.143460	0.092827	0.118143	0.088608	0.059072	0.050633
12	1.000000	0.892473	0.403226	0.000000	0.190860	0.314516	0.370968	0.368280	0.395160
13	0.969453	0.959807	0.403537	0.000000	0.205788	0.336013	0.316720	0.300643	0.310280
14	1.000000	0.409302	0.000000	0.106977	0.083721	0.088372	0.097674	0.097674	0.088372
15	1.000000	0.801252	0.425665	0.142410	0.000000	0.065728	0.143975	0.139280	0.117370
16	1.000000	0.916493	0.538622	0.089770	0.000000	0.167015	0.248434	0.229645	0.215030
17	0.841492	0.784382	0.712121	0.651515	0.579254	0.489510	0.412587	0.313520	0.228430
18	0.930124	0.934783	0.388199	0.017081	0.250000	0.333851	0.343168	0.336957	0.332290
19	0.864564	0.923933	0.465677	0.089054	0.031540	0.005566	0.000000	0.024119	0.081630

20 rows × 188 columns



In [12]: `dfs = dfs.rename(columns={187: 'classes'})`

```
In [13]: # describe the dfs dataframe in terms of metrics like count, mean, standard deviation, etc.
dfs.describe()
```

Out[13]:

	0	1	2	3	4	
count	109446.000000	109446.000000	109446.000000	109446.000000	109446.000000	109446
mean	0.891170	0.758909	0.424503	0.219602	0.201237	0
std	0.239657	0.221190	0.227561	0.207248	0.177191	0
min	0.000000	0.000000	0.000000	0.000000	0.000000	0
25%	0.922252	0.682648	0.251014	0.048853	0.082418	0
50%	0.991202	0.826560	0.430174	0.166355	0.147842	0
75%	1.000000	0.910868	0.579832	0.342707	0.259045	0
max	1.000000	1.000000	1.000000	1.000000	1.000000	1

8 rows × 188 columns



```
In [14]: #convert the last column to integer
dfs['classes']=dfs['classes'].astype(int)
```

```
In [15]: # get the number of unique values in the last column and use it to plot a bar chart
# to visualise the categories
categories =dfs['classes'].value_counts()
categories
```

Out[15]:

0	90589
4	8039
2	7236
1	2779
3	803

Name: classes, dtype: int64


In []:

```
In [17]: ▶ # Correlation matrix
def plotCorrelationMatrix(df, graphWidth):
    filename = 'ECG data'
    df = df.dropna('columns') # drop columns with NaN
    df = df[[col for col in df if df[col].nunique() > 1]] # keep columns w
    if df.shape[1] < 2:
        print(f'No correlation plots shown: The number of non-NaN or const
        return
    corr = df.corr()
    plt.figure(num=None, figsize=(graphWidth, graphWidth), dpi=80, facecol
    corrMat = plt.matshow(corr, fignum = 1)
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.columns)), corr.columns)
    plt.gca().xaxis.tick_bottom()
    plt.colorbar(corrMat)
    plt.title(f'Correlation Matrix for {filename}', fontsize=15)
    plt.show()
```

```
In [18]: plotCorrelationMatrix(dfs, 25)
```

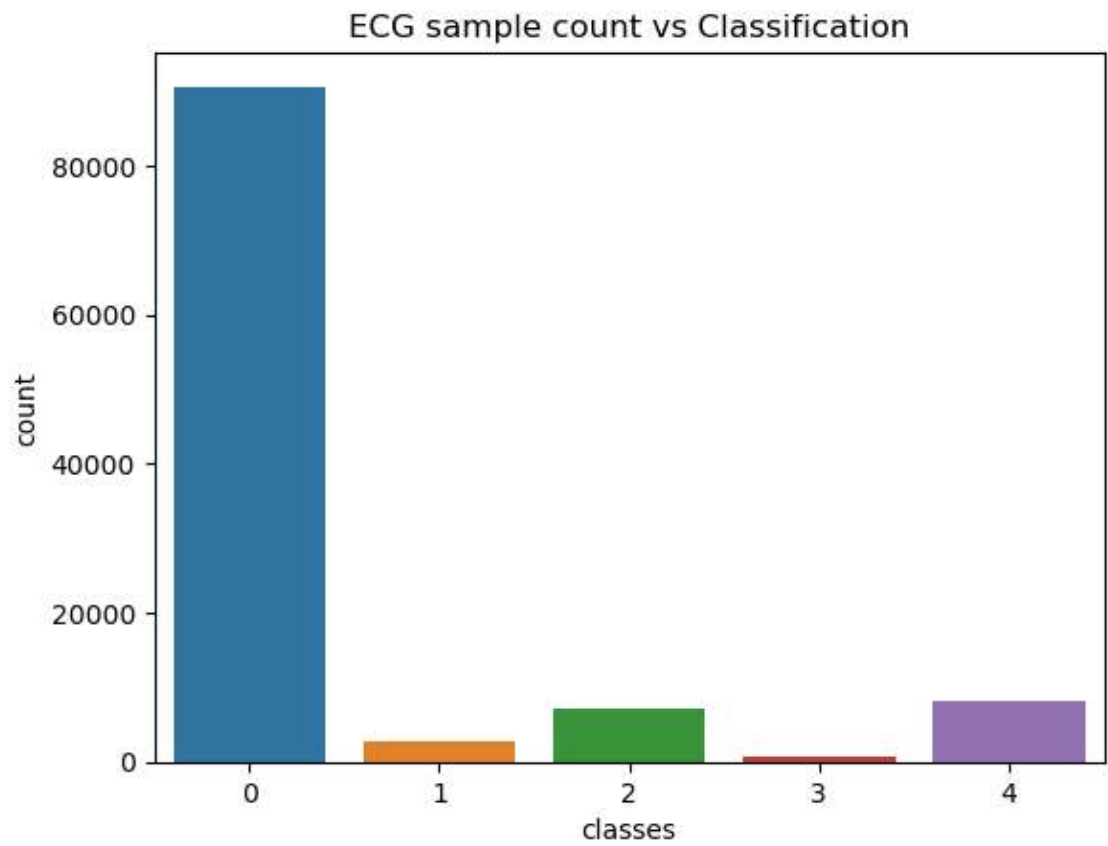
C:\Users\HP\AppData\Local\Temp\ipykernel_12656\1307321360.py:4: FutureWarning: In a future version of pandas all arguments of DataFrame.dropna will be keyword-only.

```
df = df.dropna('columns') # drop columns with NaN
```



```
In [ ]:
```

```
In [19]: ▶ # performing a countplot to visualise the data based on the number of clas
sns.countplot(x='classes', data=dfs)
plt.title('ECG sample count vs Classification')
plt.show()
```




```
In [20]: ▶ #use categories to plot a bar chart to visualise the distribution categori
# Classes: ['N': 0, 'S': 1, 'V': 2, 'F': 3, 'Q': 4]


# Create an array with the positions of the bars on the x-axis
positions = np.arange(len(categories))

# Create the bar chart
for i in categories:
    i = categories / len(dfs) * 100

plt.figure(figsize=(12,6))
plt.bar(positions, categories, color=['red','green','blue','skyblue','orange'])
plt.title('ECG vs Classification with percentage')
plt.xlabel('Classes')
plt.ylabel('Number of Samples')

# Create names on the x-axis
plt.xticks(positions, [f'N\n{i.values[0]:.2f}%', f'Q\n{i.values[1]:.2f}%',
                        f'V\n{i.values[2]:.2f}%', f'S\n{i.values[3]:.2f}%',
                        f'F\n{i.values[4]:.2f}%'])

# Show the graph
plt.show()
```



the graph shows that the category 'N' which means normal beats is in the overwhelming majority

In []: ▶

define a function to choose one sample per class

```
In [21]: ▶ def indi_sample(dfs):  
          return dfs.sample(1)  
          sample_per_class = dfs.groupby('classes', group_keys = False).apply(indi_s
```

```
In [22]: ▶ sample_per_class
```

Out[22]:

	0	1	2	3	4	5	6	7	
65298	0.824701	0.601594	0.374502	0.119522	0.000000	0.000000	0.087649	0.095618	0.12
6956	1.000000	0.882096	0.253275	0.021834	0.065502	0.126638	0.109170	0.157205	0.21
68100	0.000000	0.038462	0.108392	0.178322	0.269231	0.356643	0.444056	0.510490	0.56
65635	1.000000	0.837545	0.368231	0.252708	0.169675	0.086643	0.090253	0.093863	0.07
15414	0.496259	0.416459	0.379052	0.376559	0.336658	0.311721	0.264339	0.224439	0.17

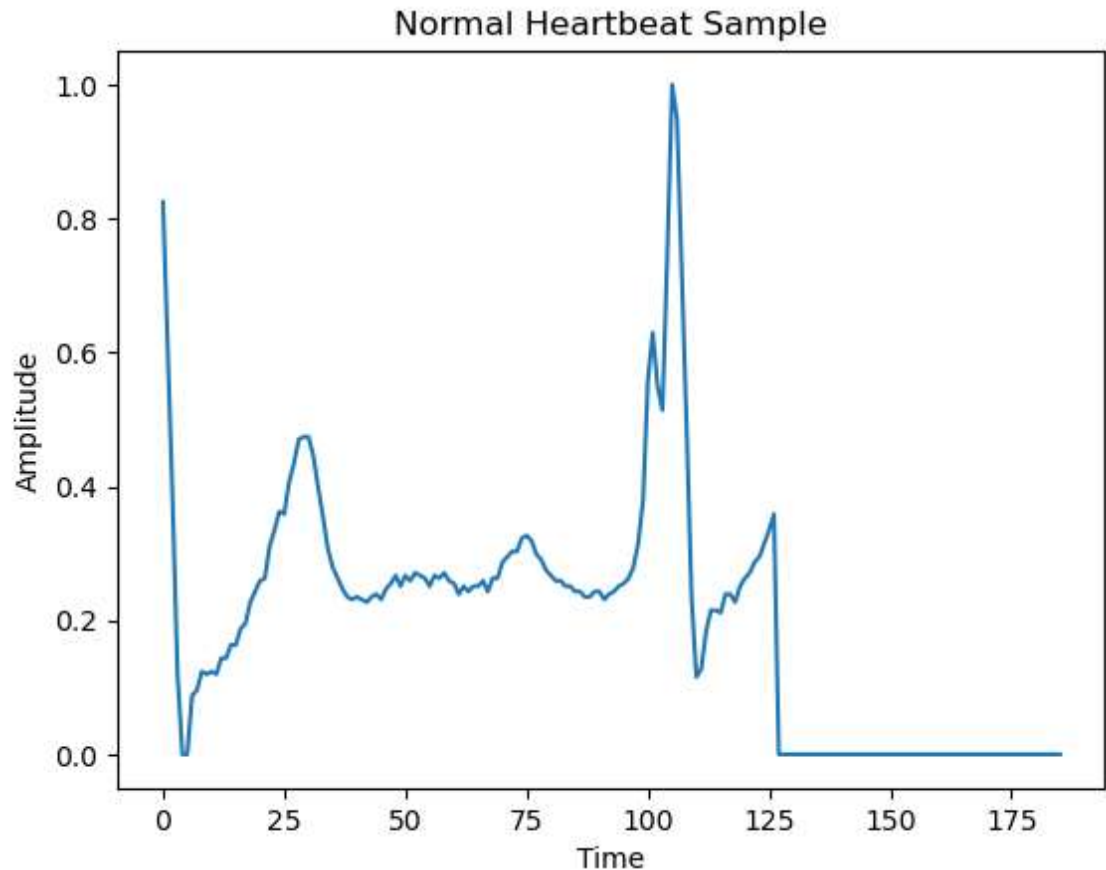
5 rows × 188 columns



from the dataframe above, we can have a look at individual examples of each ecg class


```
In [23]: ▶ # class N normal beat
plt.plot(sample_per_class.iloc[0,:186])
plt.title('Normal Heartbeat Sample')
plt.xlabel('Time')
plt.ylabel('Amplitude')
```

Out[23]: Text(0, 0.5, 'Amplitude')




```
In [24]: ▶ # class 5 supraventricular beat
plt.plot(sample_per_class.iloc[1,:186])
plt.title('supraventricular Heartbeat Sample (S)')
plt.xlabel('Time')
plt.ylabel('Amplitude')
```

Out[24]: Text(0, 0.5, 'Amplitude')



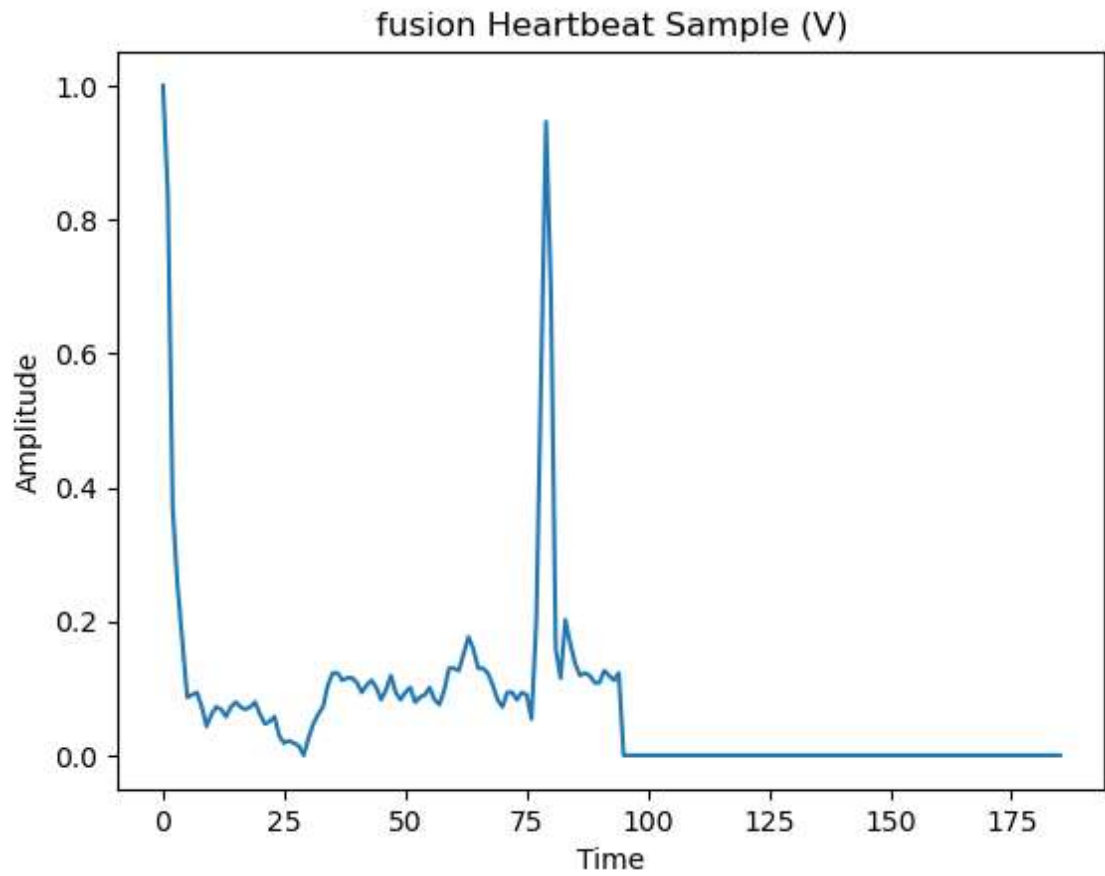
```
In [25]: ▶ # class V ventricular beat
plt.plot(sample_per_class.iloc[2,:186])
plt.title('ventricular Heartbeat Sample (V)')
plt.xlabel('Time')
plt.ylabel('Amplitude')
```

Out[25]: Text(0, 0.5, 'Amplitude')



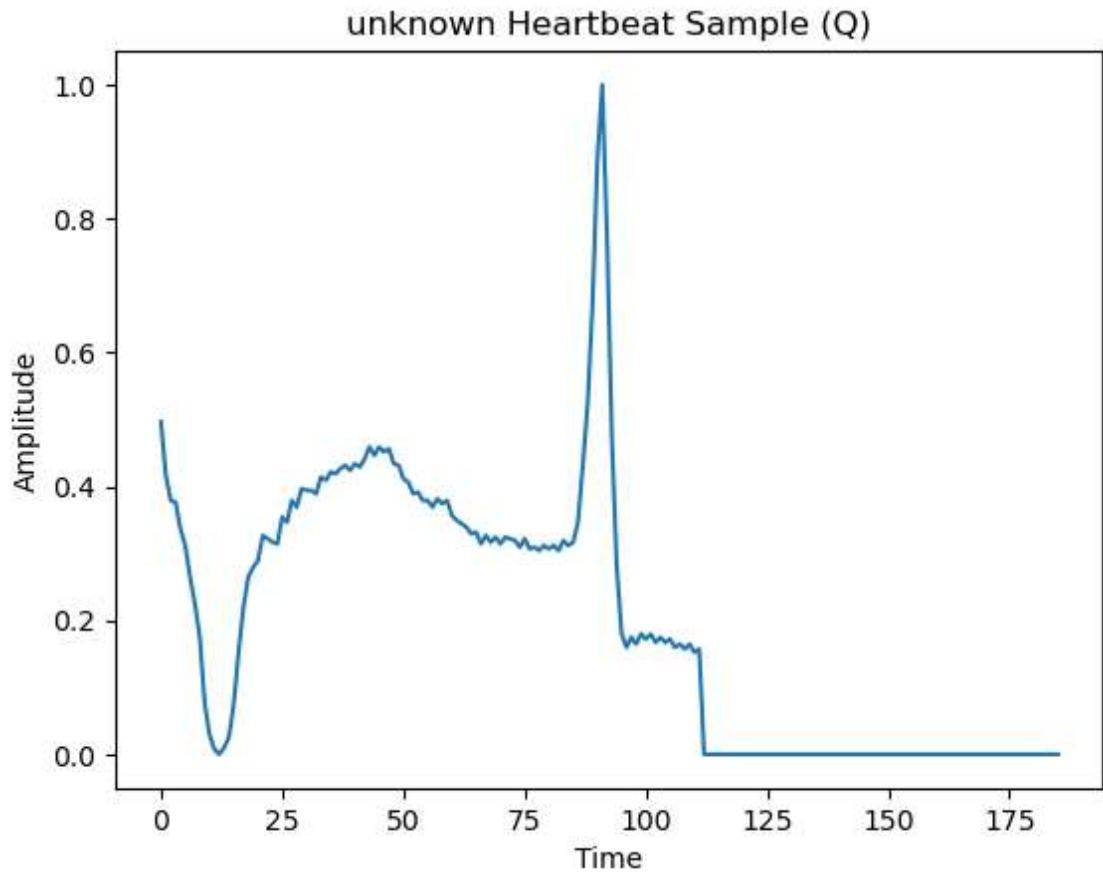
```
In [26]: ▶ # class F fusion beat
plt.plot(sample_per_class.iloc[3,:186])
plt.title('fusion Heartbeat Sample (V)')
plt.xlabel('Time')
plt.ylabel('Amplitude')
```

Out[26]: Text(0, 0.5, 'Amplitude')



```
In [27]: ▶ # class Q is unknown beat
plt.plot(sample_per_class.iloc[4,:186])
plt.title('unknown Heartbeat Sample (Q)')
plt.xlabel('Time')
plt.ylabel('Amplitude')
```

```
Out[27]: Text(0, 0.5, 'Amplitude')
```



```
In [ ]: ▶
```

Feature Selection

Here we separate the dataframe into features and target variables X and y

From observations, we know that the last column in the dataset represents the classes that signify the type of heartbeat signal. Classes: ['N': 0, 'S': 1, 'V': 2, 'F': 3, 'Q': 4]. Thus, the 'classes' column will be dropped from the dataset and used as the target or output and the rest of the dataframe, X, as the input variable.

```
In [28]: # differentiate the X inputs from the y outputs
y = dfs['classes']
```

```
In [29]: y.shape
```

```
Out[29]: (109446,)
```

```
In [30]: y
```

```
Out[30]: 0      0
         1      0
         2      0
         3      0
         4      0
         ..
109441   0
109442   0
109443   0
109444   0
109445   0
Name: classes, Length: 109446, dtype: int32
```

```
In [31]: X = dfs.drop('classes', axis = 1)
```

```
In [32]: X.shape
```

```
Out[32]: (109446, 187)
```

```
In [33]: X
```

```
Out[33]:
```

	0	1	2	3	4	5	6	7	
0	1.000000	0.691667	0.166667	0.013889	0.038889	0.036111	0.005556	0.002778	0.0
1	1.000000	0.774487	0.232346	0.143508	0.129841	0.104784	0.077449	0.091116	0.0
2	1.000000	0.892982	0.428070	0.128070	0.182456	0.198246	0.157895	0.154386	0.1
3	1.000000	0.831099	0.310992	0.040214	0.040214	0.053619	0.029491	0.008043	0.0
4	1.000000	0.907692	0.712821	0.533333	0.307692	0.153846	0.123077	0.174359	0.1
...
109441	0.916107	1.000000	0.640940	0.191275	0.083893	0.063758	0.057047	0.062081	0.0
109442	1.000000	0.830084	0.487465	0.069638	0.064067	0.203343	0.300836	0.364902	0.3
109443	0.992110	0.865878	0.299803	0.076923	0.029586	0.000000	0.007890	0.047337	0.1
109444	0.131105	0.205656	0.352185	0.465296	0.542416	0.652956	0.745501	0.781491	0.7
109445	0.980723	0.949398	0.539759	0.127711	0.065060	0.081928	0.084337	0.096386	0.0

109446 rows × 187 columns

Model Validation using RandomForest Classifier

```
In [86]: # Split the dataset into a training set and a val set
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, ran
```

```
In [87]: # model validation using the random forrest clasifier
clf = RandomForestClassifier()
```

```
In [88]: clf.fit(X_train, y_train)
```

```
Out[88]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [89]: y_validation_prediction = clf.predict(X_val)
```

```
In [90]: acc = accuracy_score (y_val, y_validation_prediction) * 100
```

```
In [91]: print(f' Validation Accuracy: {acc :.2f}%')
```

Validation Accuracy: 97.39%

```
In [92]: train_pred_RF = clf.predict(X_train)
```

```
In [ ]:
```

```
In [93]: recall_RF = recall_score(y_val, y_validation_prediction, average = 'micro')
print(f' Validation Recall: {recall_RF :.2f}%')
```

Validation Recall: 97.39%

```
In [94]: precision_RF = precision_score(y_val, y_validation_prediction, average = '
print(f' Validation precision: {precision_RF :.2f}%')
```

Validation precision: 97.39%

```
In [95]: F_score_RF = f1_score(y_val, y_validation_prediction, average = 'micro' )
print(f' Validation F-score: {F_score_RF :.2f}%')
```


Validation F-score: 97.39%

```
In [96]: ► # Generate a classification report
report = classification_report(y_val, y_validation_prediction)
print(report)
```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	18042
1	0.97	0.67	0.79	581
2	0.97	0.87	0.92	1466
3	0.89	0.54	0.67	167
4	1.00	0.95	0.97	1634
accuracy			0.97	21890
macro avg	0.96	0.80	0.87	21890
weighted avg	0.97	0.97	0.97	21890

```
In [97]: ► # Plot the confusion matrix
conmat = confusion_matrix(y_val, y_validation_prediction, labels=clf.class
display = ConfusionMatrixDisplay(confusion_matrix=conmat, display_labels=c
display.plot())
```

Out[97]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20501a43190>



MLP

```
In [98]: ▶ # Split the dataset into a training set and a test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [99]: ▶ clfMLP = MLPClassifier(hidden_layer_sizes=(100, 500,300,500, 100), activat
```

```
In [100]: ▶ clfMLP.fit(X_train, y_train)
```

```
Out[100]: ▼ MLPClassifier
MLPClassifier(alpha=0.25, hidden_layer_sizes=(100, 500, 300, 500, 100),
max_iter=2000)
```

```
In [ ]: ▶
```

```
In [101]: ▶ # Make predictions on the test data
y_prediction = clfMLP.predict(X_test)
```

```
In [102]: ▶ train_pred = clfMLP.predict(X_train)
```

```
In [103]: ▶ train_acc_MLP = accuracy_score( y_train, train_pred) * 100
print(f'train Accuracy: {train_acc_MLP:.2f}%')
```

```
train Accuracy: 98.28%
```

```
In [ ]: ▶
```

```
In [104]: ▶ accMLP = accuracy_score(y_test, y_prediction) * 100
print(f'test Accuracy: {accMLP :.2f}%')
```

```
test Accuracy: 97.80%
```

```
In [105]: ▶ recall_MLP = recall_score(y_test, y_prediction, average = 'micro') * 100
print(f' MLP Recall: {recall_MLP :.2f}%')

precision_MLP = precision_score(y_test, y_prediction, average = 'micro' )
print(f' MLP precision: {precision_MLP :.2f}%')

F_score_MLP = f1_score(y_test, y_prediction, average = 'micro' ) * 100
print(f' MLP F-score: {F_score_MLP :.2f}%')
```


```
MLP Recall: 97.80%
MLP precision: 97.80%
MLP F-score: 97.80%
```

```
In [106]: ► # Generate a classification report
report_MLP = classification_report(y_test, y_prediction)
print(report_MLP)
```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	18092
1	0.92	0.70	0.79	560
2	0.95	0.93	0.94	1486
3	0.70	0.63	0.66	156
4	0.99	0.97	0.98	1596
accuracy			0.98	21890
macro avg	0.91	0.84	0.87	21890
weighted avg	0.98	0.98	0.98	21890

```
In [107]: ► # Plot the confusion matrix
conmat_MLP = confusion_matrix(y_test, y_prediction, labels=clfMLP.classes_)
display = ConfusionMatrixDisplay(confusion_matrix=conmat_MLP, display_labels=clfMLP.classes_)
display.plot()
```

Out[107]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20512db4d90>



SVM

```
In [57]:  from sklearn.svm import SVC
```

```
In [58]:  clf_SVC = SVC(kernel = 'poly')
```

```
In [59]:  clf_SVC.fit(X_train, y_train)
```

```
Out[59]:  SVC
          SVC(kernel='poly')
```

```
In [60]:  y_pred = clf_SVC.predict(X_test)
          y_pred
```

```
Out[60]:  array([0, 0, 0, ..., 2, 2, 0])
```

```
In [61]:  training_pred = clf_SVC.predict(X_train)
```

```
In [62]:  training_accuracy_SVC = accuracy_score( y_train, training_pred) * 100
          print(f'train Accuracy: {training_accuracy_SVC:.2f}%')
```

```
train Accuracy: 96.97%
```

```
In [64]:  acc_SVC = accuracy_score(y_test, y_pred) * 100
          print(f'test Accuracy: {acc_SVC :.2f}%')
```

```
test Accuracy: 96.53%
```

```
In [65]:  recall_SVC = recall_score(y_test, y_pred, average = 'micro') * 100
          print(f' MLP Recall: {recall_SVC :.2f}%')

          precision_SVC = precision_score(y_test, y_pred, average = 'micro' ) * 100
          print(f' MLP precision: {precision_SVC :.2f}%')

          F_score_SVC = f1_score(y_test, y_pred, average = 'micro' ) * 100
          print(f' MLP F-score: {F_score_SVC :.2f}%')
```


```
MLP Recall: 96.53%
MLP precision: 96.53%
MLP F-score: 96.53%
```

```
In [66]: # Generate a classification report
report_SVC = classification_report(y_test, y_pred)
print(report_SVC)
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	18092
1	0.97	0.59	0.73	560
2	0.96	0.85	0.90	1486
3	0.77	0.40	0.53	156
4	0.99	0.90	0.94	1596
accuracy			0.97	21890
macro avg	0.93	0.75	0.82	21890
weighted avg	0.96	0.97	0.96	21890

```
In [67]: # Plot the confusion matrix
conmat_SVC = confusion_matrix(y_test, y_pred, labels=clf_SVC.classes_)
display = ConfusionMatrixDisplay(confusion_matrix=conmat_SVC, display_labels=clf_SVC.classes_)
display.plot()
```

```
Out[67]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20501790d90>
```



Comparing the performance metrics on all the models

```
In [108]: ▶ #table of metrics
model_eval= {'models':['MLP', 'SVM', 'RFS'],
'Accuracies':[accMLP, acc_SVC, acc],
'precisions':[precision_MLP, precision_SVC, precision_RF],
'recalls': [recall_MLP, recall_SVC, recall_RF],
'f_scores' : [F_score_MLP, F_score_SVC, F_score_RF]
}

model_table = pd.DataFrame(model_eval)
model_table
```


Out[108]:

	models	Accuracies	precisions	recalls	f_scores
0	MLP	97.802650	97.802650	97.802650	97.802650
1	SVM	96.528095	96.528095	96.528095	96.528095
2	RFS	97.386935	97.386935	97.386935	97.386935

Accuracy

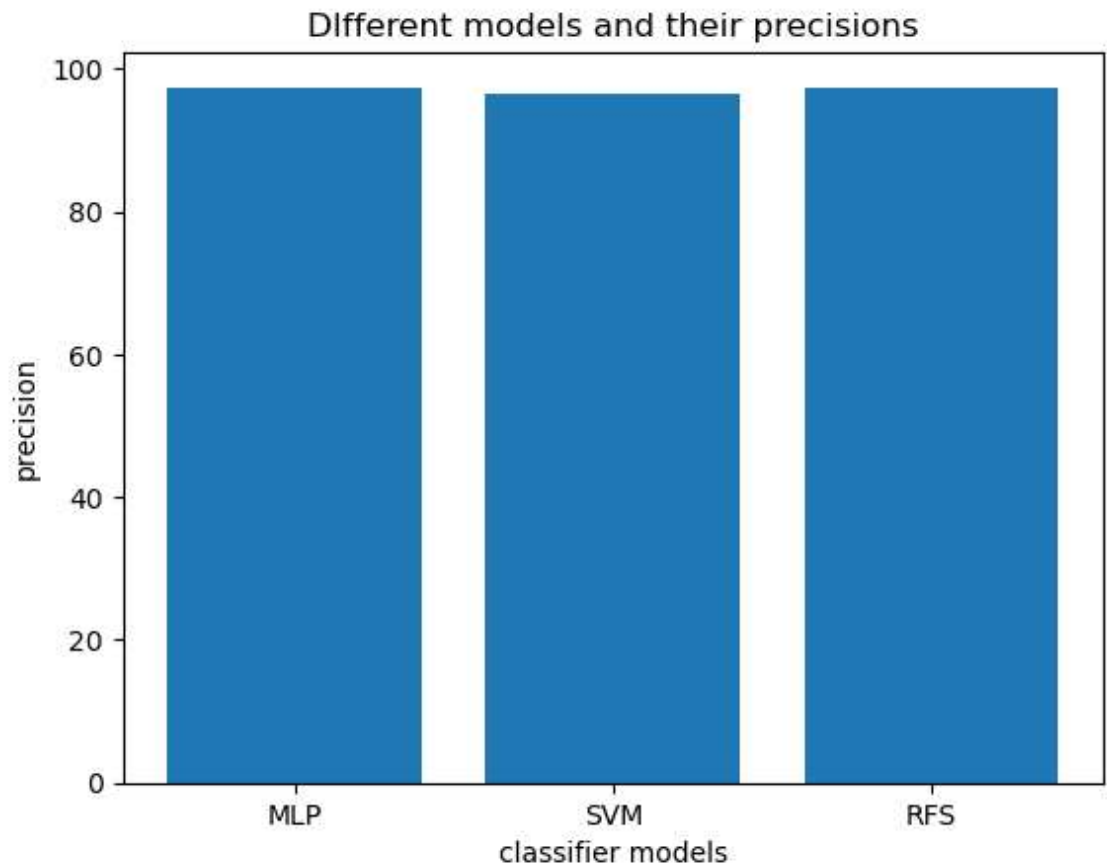
```
In [70]: ▶ testAcc = [accMLP, acc_SVC, acc]
precisions =[precision_MLP, precision_SVC, precision_RF]
recalls= [recall_MLP, recall_SVC, recall_RF]
f_scores = [F_score_MLP, F_score_SVC, F_score_RF]
```

```
In [71]: ▶ models = ['MLP', 'SVM', 'RFS']  
plt.bar(models, testAcc)  
plt.xlabel('classifier models')  
plt.ylabel('accuracy')  
plt.title('Different models and their accuracies')  
plt.show()
```




Precision

```
In [80]: ▶ models = ['MLP', 'SVM', 'RFS']  
plt.bar(models, precisions)  
plt.xlabel('classifier models')  
plt.ylabel('precision')  
plt.title('Different models and their precisions')  
plt.show()
```



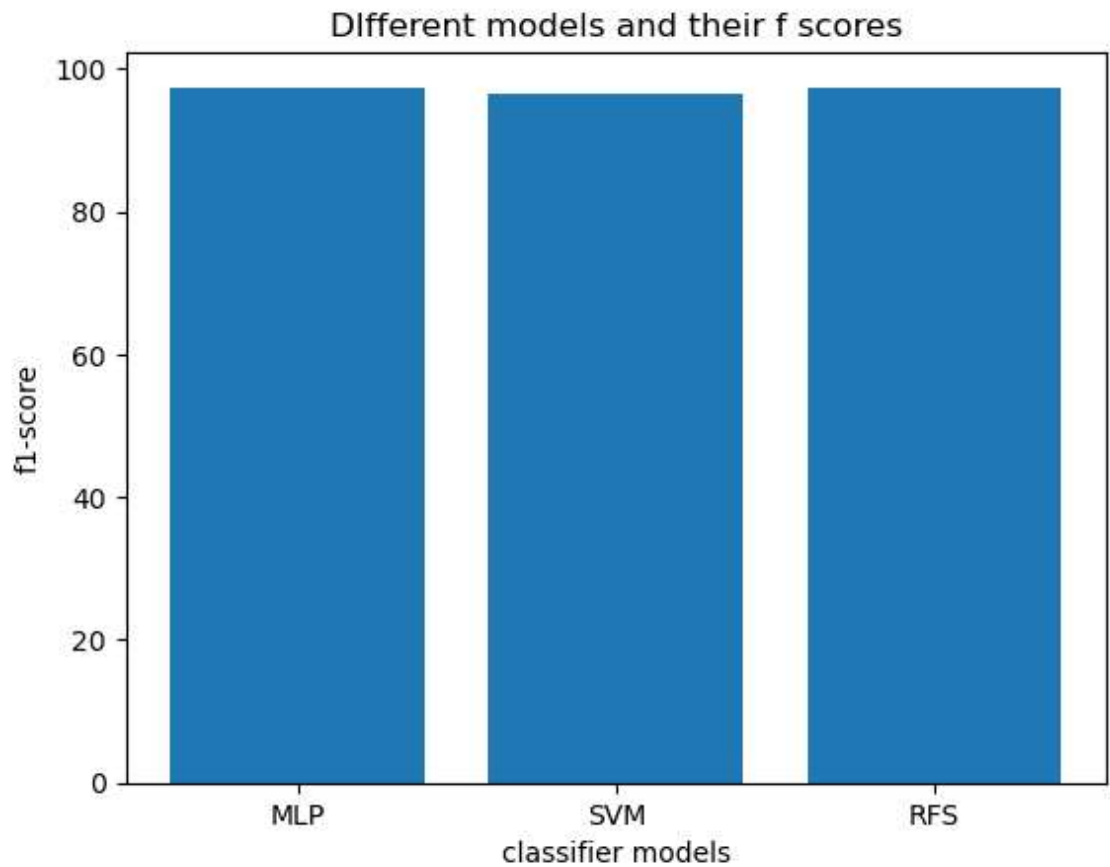
Recall

```
In [81]: ▶ models = ['MLP', 'SVM', 'RFS']  
plt.bar(models, recalls)  
plt.xlabel('classifier models')  
plt.ylabel('recall')  
plt.title('Different models and their recalls')  
plt.show()
```



F-score

```
In [82]: ▶ models = ['MLP', 'SVM', 'RFS']  
plt.bar(models, f_scores)  
plt.xlabel('classifier models')  
plt.ylabel('f1-score')  
plt.title('Different models and their f scores')  
plt.show()
```



```
In [ ]: ▶
```